

# Programovanie v oeračných systémoch

## 03 - Filesystem IO

Jozef Šiška



Department of Applied Informatics  
Comenius University in Bratislava

2016/2017



- VFS (virtual filesystem)
- mounted "real" filesystems (`mount(2)`)
- files
  - name, data, metadata
  - inode
  - `open (creat)`, `close`, `read`, `write`, `stat`, ...
- directories ("folders")
  - list of entries (files, directories)
  - `open`, `close`, `readdir`, `getdents`, ...

# Filesystem related calls

`fd = open(path, flags)`

`close(fd)`

`count = read(fd, &buf, count)`

`count = write(fd, &buf, count)`

`pos = lseek(fd, offset, whence)`

Open/create a file

Close a file

Read from file

Write to a file

Change position

`stat(path, &buf), fstat(fd, &buf)`

`access(name, mode)`

`chmod(n, mode), chown(n, u, g)`

`utime(name, times)`

`umask(mode)`

Get file (status) information

Check accessibility / existence

Change permissions, owner

Change file access / modification times

Change file creation mode mask

`rename(old, new)`

`mkdir(name, mode), rmdir(name)`

`link(n1, n2), unlink(name)`

Rename file

Create / remove empty directory

Create link, remove dir entry

`mount(special, name, flags)`

`umount(special)`

`sync(), syncfs(), fsync(), fdatasync()`

Mount fs

Unmount fs

Synchronize file(system) to disk

`chdir(dirname), chroot(dirname)`

Change current / root directory

## ... and more

<code>dup(fd), dup2(fd,newfd), dup3(..., flags)</code>	Duplicate fd
<code>fcntl(fd, cmd, args...)</code>	File "operations"
<code>ioctl(fd, request, args...)</code>	"Special" operations on fd

Maybe not a file:

- Regular files, directories
- Block and character devices
- Pipes
- Sockets

<code>pipe, pipe2</code>	Creates a pipe
<code>mknod</code>	Create "special" (device) files

## ... and more

<code>dup(fd), dup2(fd,newfd), dup3(..., flags)</code>	Duplicate fd
<code>fcntl(fd, cmd, args...)</code>	File "operations"
<code>ioctl(fd, request, args...)</code>	"Special" operations on fd

Maybe not a file:

- Regular files, directories
- Block and character devices
- Pipes
- Sockets

<code>pipe, pipe2</code>	Creates a pipe
<code>mknod</code>	Create "special" (device) files

# Non-blocking IO

- `open(..., O_NONBLOCK)`
- `fcntl` on already open fds

```
int flags = fcntl(fd, F_GETFL, 0); // check for -1
fcntl(fd, F_SETFL, flags | O_NONBLOCK); // check for -1
```

- read/write: return `EAGAIN`, `EWOULDBLOCK`
- `select`, `poll`, `epoll`: Wait for events on file descriptors
- `ioctl(fd, FIONREAD, &bytes_available)`

Filesystem io always returns the full file size as available, and will block (if data needs to be read from disk etc.).

# Get output of a program

```
int pipefd[2];
pid_t pid;

pipe(pipefd);

pid = fork();
if (pid == 0) { // child
    dup2(pipefd[1], 1);
    execvp(argv[1], argv+1);
} else { // parent
    char buffer[1024];
    ssize_t br = read(pipefd[0], buffer, sizeof(buffer));
    printf("Output: %.*s\n", (int)br, buffer);
    int wstatus;
    if (waitpid(pid, &wstatus, 0) == -1) {
        perror("waitpid");
        exit(EXIT_FAILURE);
    }
    if (WIFEXITED(wstatus)) {
        printf("Exited: %d\n", WEXITSTATUS(wstatus));
    }
}
return 0;
```